

Resilient Distributed Datasets: A Fault-Tolerant Abstraction for In-Memory Cluster Computing

NOTE: Your slides/presentation need to cover the assigned sections and questions in a clear and well-organized manner. You are allowed to borrow contents from other resources, such as online slides, as long as you acknowledge them. For a slide that covers a given question, please print the question on the slide. However, you don't have to answer the question using a long paragraph of text on the slides. Instead, use bullet points, graph, animation, or oral explanation to answer the question. In your Q&A report, use text to more thoroughly answer the questions.

You only need to cover contents until Section 5.1.

- (1) “...individual RDDs are immutable...” What does it mean by being “immutable”? What benefits does this property of RDD bring?
- (2) When an RDD is being created (new data are being written into it), can the data in the RDD be read for computing before the RDD is completed created?
- (3) “This allows them to efficiently provide fault tolerance by logging the transformations used to build a dataset (its lineage) rather than the actual data.” “To achieve fault tolerance efficiently, RDDs provide a restricted form of shared memory, based on coarse-grained transformations rather than fine-grained updates to shared state.” Why does using RDD help to provide efficient fault tolerance? Or why does coarse-grained transformation help with the efficiency?
- (4) What is difference between transformation and action?
- (5) “In addition, programmers can call a persist method to indicate which RDDs they want to reuse in future operations.” What's the consequence if a user does not explicitly request persistence of an RDD?
- (6) Explain Figure 1 about a lineage graph.